

```
In [1]: # Step 1: Load & Inspect Dataset

import pandas as pd

# Load the dataset
df = pd.read_csv('/Users/maazhussain/Desktop/Projects/Banking Churn Risk & F

# Show structure and first rows
print("Shape:", df.shape)
df.info()
df.head()
```

```
Shape: (10000, 14)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary         10000 non-null  float64
13  Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
Out[1]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1
2	3	15619304	Onio	502	France	Female	42	8
3	4	15701354	Boni	699	France	Female	39	1
4	5	15737888	Mitchell	850	Spain	Female	43	2

```
In [3]: # Step 2: Clean & Prepare

# Drop non-essential columns
df_clean = df.drop(columns=["RowNumber", "CustomerId", "Surname"])

# Check the cleaned dataset
print("Cleaned shape:", df_clean.shape)
df_clean.head()
```

Cleaned shape: (10000, 11)

Out[3]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCr
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	
4	850	Spain	Female	43	2	125510.82	1	

In [5]: *# Step 3: EDA – Visualize churn vs key features*

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.set(style="whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

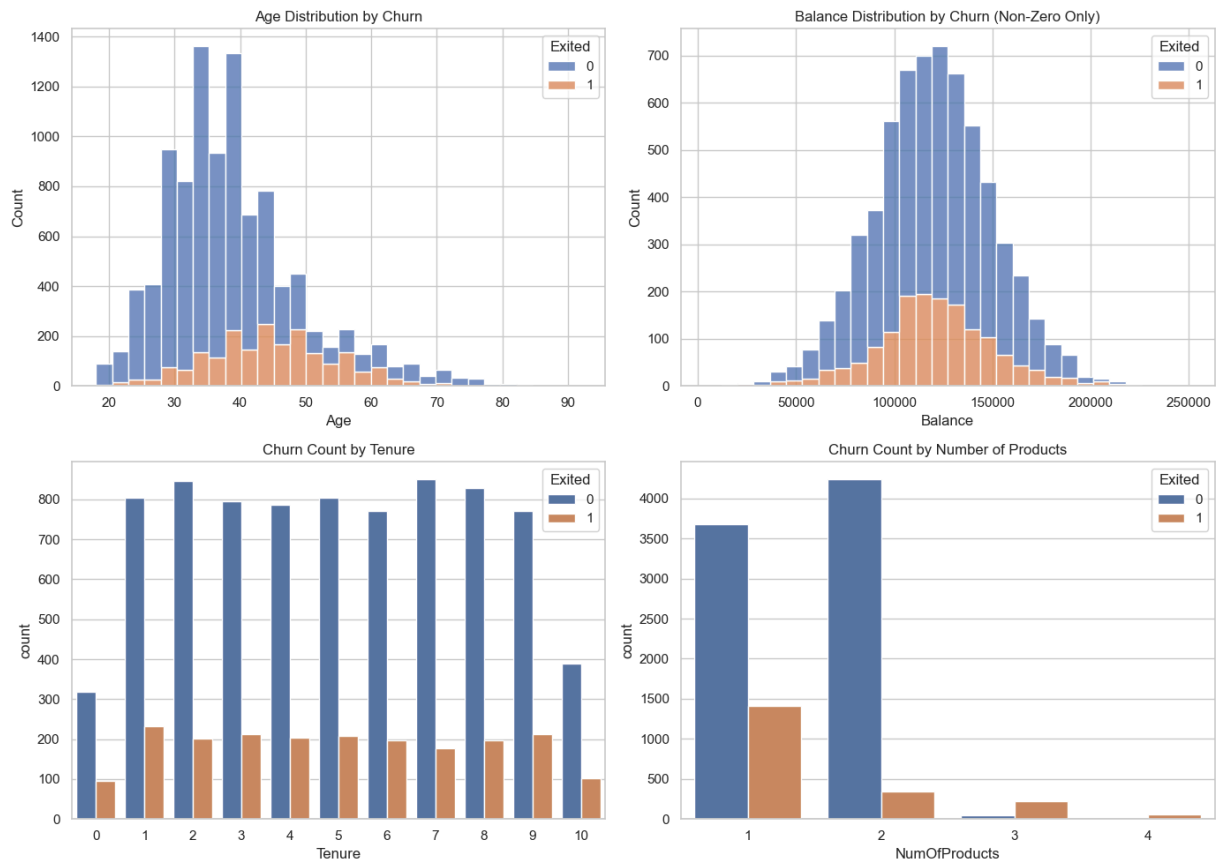
# 1. Age vs Churn
sns.histplot(data=df_clean, x="Age", hue="Exited", multiple="stack", bins=30)
axes[0, 0].set_title("Age Distribution by Churn")

# 2. Balance vs Churn (exclude zero balances)
sns.histplot(data=df_clean[df_clean["Balance"] > 0], x="Balance", hue="Exited", multiple="stack")
axes[0, 1].set_title("Balance Distribution by Churn (Non-Zero Only)")

# 3. Tenure vs Churn
sns.countplot(data=df_clean, x="Tenure", hue="Exited", ax=axes[1, 0])
axes[1, 0].set_title("Churn Count by Tenure")

# 4. Num of Products vs Churn
sns.countplot(data=df_clean, x="NumOfProducts", hue="Exited", ax=axes[1, 1])
axes[1, 1].set_title("Churn Count by Number of Products")

plt.tight_layout()
plt.show()
```



In [7]: *# Step 4: Churn Breakdown by Categorical Features*

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.set(style="whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

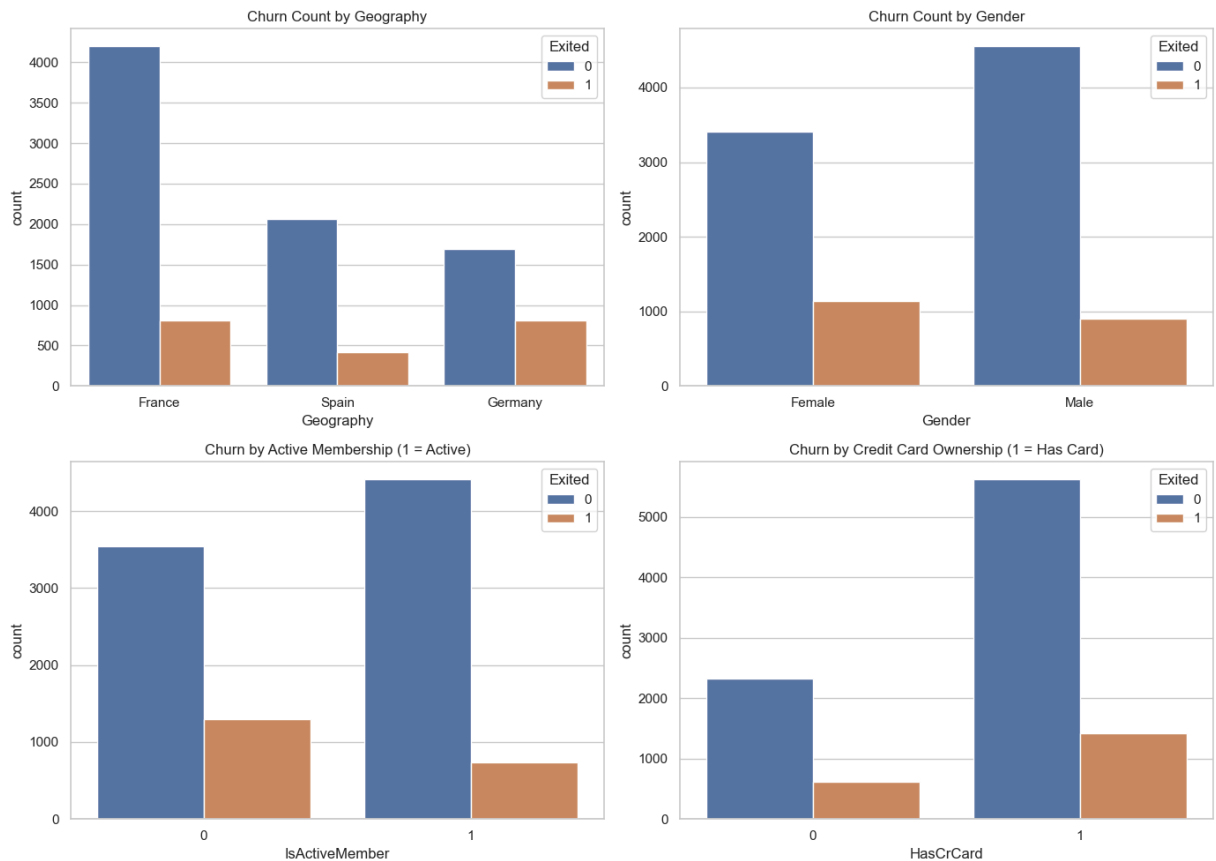
# 1. Geography vs Churn
sns.countplot(data=df_clean, x="Geography", hue="Exited", ax=axes[0, 0])
axes[0, 0].set_title("Churn Count by Geography")

# 2. Gender vs Churn
sns.countplot(data=df_clean, x="Gender", hue="Exited", ax=axes[0, 1])
axes[0, 1].set_title("Churn Count by Gender")

# 3. IsActiveMember vs Churn
sns.countplot(data=df_clean, x="IsActiveMember", hue="Exited", ax=axes[1, 0])
axes[1, 0].set_title("Churn by Active Membership (1 = Active)")

# 4. HasCrCard vs Churn
sns.countplot(data=df_clean, x="HasCrCard", hue="Exited", ax=axes[1, 1])
axes[1, 1].set_title("Churn by Credit Card Ownership (1 = Has Card)")

plt.tight_layout()
plt.show()
```



In [9]: *# Step 5.1: Prepare data for modeling*

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Copy the clean dataframe
model_df = df_clean.copy()

# Encode categorical columns: Geography and Gender
model_df["Geography"] = LabelEncoder().fit_transform(model_df["Geography"])
model_df["Gender"] = LabelEncoder().fit_transform(model_df["Gender"])

# Define features and target
X = model_df.drop("Exited", axis=1)
y = model_df["Exited"]

# Split into training and testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X_train.shape, X_test.shape
```

Out[9]: ((8000, 10), (2000, 10))

In [11]: *# Step 5.2: Train & Evaluate Logistic Regression*

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Initialize and train model
```

```

lr = LogisticRegression(max_iter=1000)
lr.fit(X_train, y_train)

# Predict on test set
y_pred = lr.predict(X_test)

# Evaluation
print(" Accuracy:", accuracy_score(y_test, y_pred))
print("\n Classification Report:\n", classification_report(y_test, y_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Accuracy: 0.8165

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.97	0.89	1607
1	0.61	0.18	0.28	393
accuracy			0.82	2000
macro avg	0.72	0.58	0.59	2000
weighted avg	0.79	0.82	0.77	2000

Confusion Matrix:

```

[[1561  46]
 [ 321  72]]

```

/opt/anaconda3/lib/python3.12/site-packages/sklearn/linear\_model/\_logistic.p  
y:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

In [13]: *# Step 5.3: Random Forest Classifier*

```

from sklearn.ensemble import RandomForestClassifier

# Initialize and train
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Predict
rf_pred = rf.predict(X_test)

# Evaluation
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

print(" Random Forest Accuracy:", accuracy_score(y_test, rf_pred))
print("\n Classification Report:\n", classification_report(y_test, rf_pred))
print("\n Confusion Matrix:\n", confusion_matrix(y_test, rf_pred))

```

Random Forest Accuracy: 0.8645

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.96	0.92	1607
1	0.75	0.47	0.58	393
accuracy			0.86	2000
macro avg	0.81	0.71	0.75	2000
weighted avg	0.85	0.86	0.85	2000

Confusion Matrix:

```
[[1545  62]
 [ 209 184]]
```

## Capstone Report – Banking Churn Risk Analysis

### \*Objective\*

To analyze customer churn in a European bank and identify key drivers using exploratory data analysis and predictive modeling. The goal is to support data-driven retention strategies.

### \*Dataset Summary\*

Source: Kaggle – Churn for Bank Customers Rows: 10,000 customers Target Variable: Exited (1 = churned, 0 = retained) Key Features: Age, Geography, Tenure, Balance, Products, Credit Card Ownership, Activity

### \*Key Insights from EDA\*

Feature Churn Trend Age Older customers (40+) churn more frequently Balance High churn even among high-balance customers Products Single-product customers are most likely to churn Geography Germany has the highest churn ratio among regions Activity Inactive users show significantly higher churn Credit Card Non-card holders churn slightly more often

### \*Model Results\*

Model | Accuracy | Churn Recall | Churn Precision | F1 (Churn) Logistic Regression | 81.7% | 18% | 61% | 28% Random Forest | 86.5% | 47% | 75% | 58%

**\*Conclusion:** Random Forest significantly outperforms logistic regression, capturing more than double the churners. It's recommended for deployment.

### \*Recommendations\*

Focus retention efforts on older, inactive customers with only one product Target Germany with tailored campaigns Promote cross-selling to increase product ownership

Build automated churn risk dashboards for real-time decision-making

In [ ]: